

# EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks

---

Neural Acceleration Lab

Hwigeon Oh

# Introduction

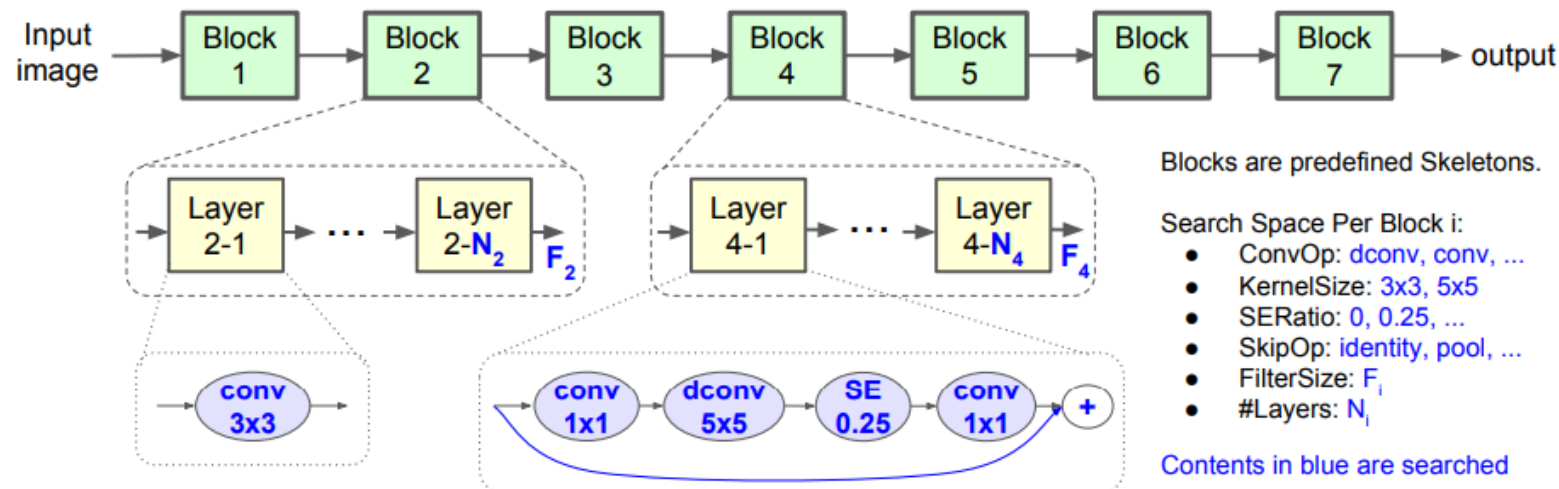
- Scaling up ConvNets is widely used to achieve better accuracy
- Ways to scale up: Depth (#layers) / Width (#channels) / Image resolution
- In previous work, it is common to scale only one of three
- This paper provides a principled method to scale up with ***compound scaling***

# Problem Formulation

- A ConvNet Layer  $i$  can be defined as:  $Y_i = \mathcal{F}_i(X_i)$
- A ConvNet  $\mathcal{N}$  can be represented as:  $\mathcal{N} = \mathcal{F}_k \odot \cdots \odot \mathcal{F}_2 \odot \mathcal{F}_1(X_1) = \odot_{j=1 \dots k} \mathcal{F}_j(X_1)$
- By factorization with blocks:  $\mathcal{N} = \odot_{i=1 \dots s} \mathcal{F}_i^{L_i}(X_{\langle H_i, W_i, C_i \rangle})$

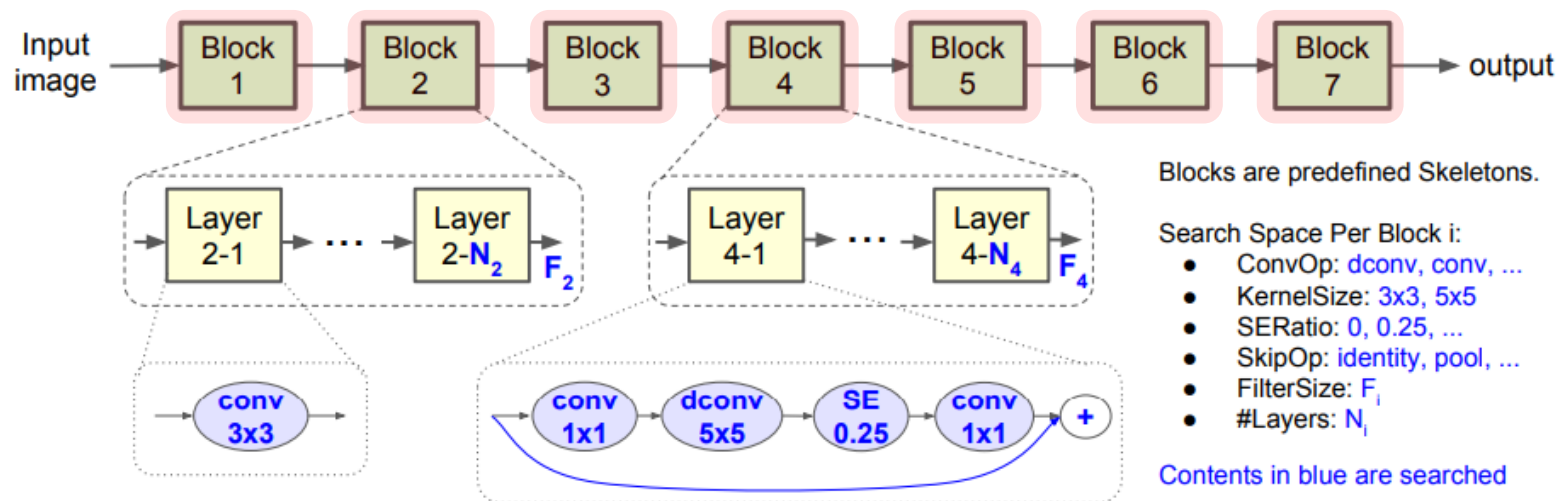
# Problem Formulation

- A ConvNet Layer  $i$  can be defined as:  $Y_i = \mathcal{F}_i(X_i)$
- A ConvNet  $\mathcal{N}$  can be represented as:  $\mathcal{N} = \mathcal{F}_k \odot \cdots \odot \mathcal{F}_2 \odot \mathcal{F}_1(X_1) = \odot_{j=1 \dots k} \mathcal{F}_j(X_1)$
- By factorization with blocks:  $\mathcal{N} = \odot_{i=1 \dots s} \mathcal{F}_i^{L_i}(X_{\langle H_i, W_i, C_i \rangle})$



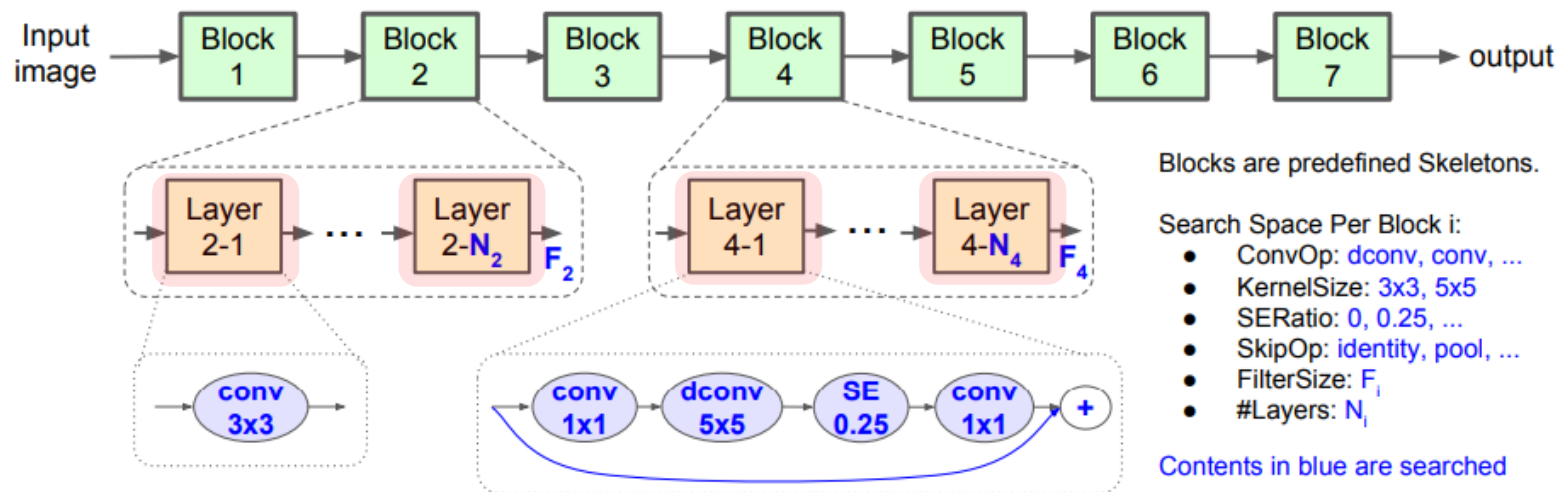
# Problem Formulation

- A ConvNet Layer  $i$  can be defined as:  $Y_i = \mathcal{F}_i(X_i)$
- A ConvNet  $\mathcal{N}$  can be represented as:  $\mathcal{N} = \mathcal{F}_k \odot \cdots \odot \mathcal{F}_2 \odot \mathcal{F}_1(X_1) = \odot_{j=1 \dots k} \mathcal{F}_j(X_1)$
- By factorization with blocks:  $\mathcal{N} = \odot_{i=1 \dots s} \mathcal{F}_i^{L_i}(X_{\langle H_i, W_i, C_i \rangle})$



# Problem Formulation

- A ConvNet Layer  $i$  can be defined as:  $Y_i = \mathcal{F}_i(X_i)$
- A ConvNet  $\mathcal{N}$  can be represented as:  $\mathcal{N} = \mathcal{F}_k \odot \cdots \odot \mathcal{F}_2 \odot \mathcal{F}_1(X_1) = \odot_{j=1 \dots k} \mathcal{F}_j(X_1)$
- By factorization with blocks:  $\mathcal{N} = \odot_{i=1 \dots s} \mathcal{F}_i^{L_i}(X_{\langle H_i, W_i, C_i \rangle})$





# Problem Formulation

- **Model scaling tries to expand** the length( $L_i$ ), width( $C_i$ ), resolution( $H_i, W_i$ )
- **Without changing layer architecture**( $\mathcal{F}_i$ )
- To reduce the design space, all layers are **scaled uniformly** with constant ratio

$$\begin{aligned} \max_{d,w,r} \quad & \text{Accuracy}(\mathcal{N}(d, w, r)) \\ \text{s.t.} \quad & \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i} (X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, w \cdot \hat{C}_i \rangle}) \\ & \text{Memory}(\mathcal{N}) \leq \text{target\_memory} \\ & \text{FLOPS}(\mathcal{N}) \leq \text{target\_flops} \end{aligned}$$

# Scaling Dimensions: depth

- Deeper network can capture **richer** and **more complex** features
- Ex. ResNet-50 < ResNet-152





$$\begin{aligned} \max_{d,w,r} \quad & \text{Accuracy}(\mathcal{N}(d, w, r)) \\ \text{s.t.} \quad & \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d, \hat{L}_i} (X_{\langle r, \hat{H}_i, r, \hat{W}_i, w, \hat{C}_i \rangle}) \\ & \text{Memory}(\mathcal{N}) \leq \text{target\_memory} \\ & \text{FLOPS}(\mathcal{N}) \leq \text{target\_flops} \end{aligned}$$



# Scaling Dimensions: width

$$\begin{aligned} \max_{d,w,r} \quad & \text{Accuracy}(\mathcal{N}(d, \boxed{w}, r)) \\ \text{s.t.} \quad & \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i} (X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, \boxed{w} \cdot \hat{C}_i \rangle}) \\ & \text{Memory}(\mathcal{N}) \leq \text{target\_memory} \\ & \text{FLOPS}(\mathcal{N}) \leq \text{target\_flops} \end{aligned}$$



- Wider network can capture **more fine-grained** features
- Ex. wide residual network



# Scaling Dimensions: resolution

$$\begin{aligned} & \max_{d,w,r} \text{Accuracy}(\mathcal{N}(d, w, r)) \\ & s.t. \quad \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i} (X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, w \cdot \hat{C}_i \rangle}) \\ & \quad \text{Memory}(\mathcal{N}) \leq \text{target\_memory} \\ & \quad \text{FLOPS}(\mathcal{N}) \leq \text{target\_flops} \end{aligned}$$




- With higher resolution input, can capture **more fine-grained** patterns
- Ex. GPipe



# Scaling Dimensions


$$\begin{aligned} & \max_{d,w,r} \text{Accuracy}(\mathcal{N}(d, w, r)) \\ & s.t. \quad \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i} (X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, w \cdot \hat{C}_i \rangle}) \\ & \quad \text{Memory}(\mathcal{N}) \leq \text{target\_memory} \\ & \quad \text{FLOPS}(\mathcal{N}) \leq \text{target\_flops} \end{aligned}$$

- Scaling up any dimension improves accuracy
- The accuracy gain diminishes for bigger models




# Compound Scaling

- Scaling dimensions are not independent (empirical observation)
- Increase depth to make larger receptive field for higher resolution images
- Increase width to capture more fine-grained patterns with more pixels



# Compound Scaling

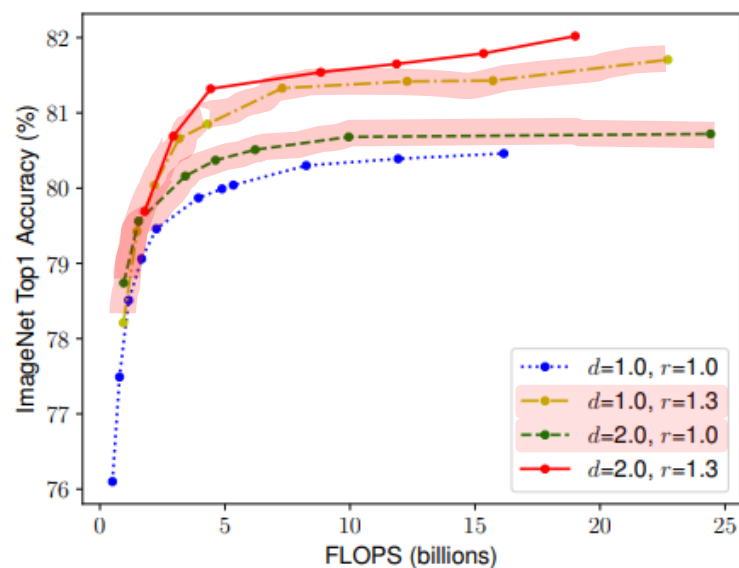
- Scaling dimensions are not independent (empirical observation)
- Increase depth to make larger receptive field for higher resolution images
- Increase width to capture more fine-grained patterns with more pixels



- Each dot in a line denotes a model with different width
- Accuracy saturates quickly if only scale width

# Compound Scaling


- Scaling dimensions are not independent (empirical observation)
- Increase depth to make larger receptive field for higher resolution images
- Increase width to capture more fine-grained patterns with more pixels



- Each dot in a line denotes a model with different width
- Accuracy saturates quickly if only scale width
- Scaling width with depth or resolution makes better

# Compound Scaling

- Scaling dimensions are not independent (empirical observation)
- Increase depth to make larger receptive field for higher resolution images
- Increase width to capture more fine-grained patterns with more pixels



- Each dot in a line denotes a model with different width
- Accuracy saturates quickly if only scale width
- Scaling width with depth or resolution makes better
- **It is critical to balance all three dimensions**

# Compound Scaling Method

- Depth:  $d = \alpha^\phi$ , Width:  $w = \beta^\phi$ , Resolution:  $r = \gamma^\phi$
- Use a compound coefficient  $\phi$  to uniformly scales
- The FLOPS of a convolution op is proportional to  $d, w^2, r^2$
- With constraint  $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$ , the total FLOPS will increase by  $2^\phi$
- $\phi$  controls how many more resources are available
- $\alpha, \beta, \gamma$  specify how to assign extra resources



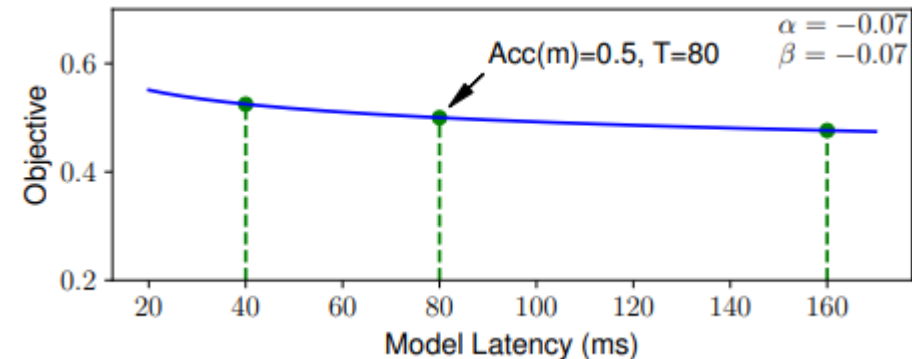
# MnasNet (Tan et al., 2019)

- Multi-objective neural architecture search that optimized both accuracy and latency
- Use weighted product method to approximate Pareto optimal solutions
- Use empirical observation: x2 the latency brings about 5% relative accuracy gain
  - With this condition,  $\beta$  in the below equation is  $\approx -0.07$

$$\begin{aligned} & \text{maximize}_m \quad ACC(m) \\ & \text{subject to} \quad LAT(m) \leq T \end{aligned}$$

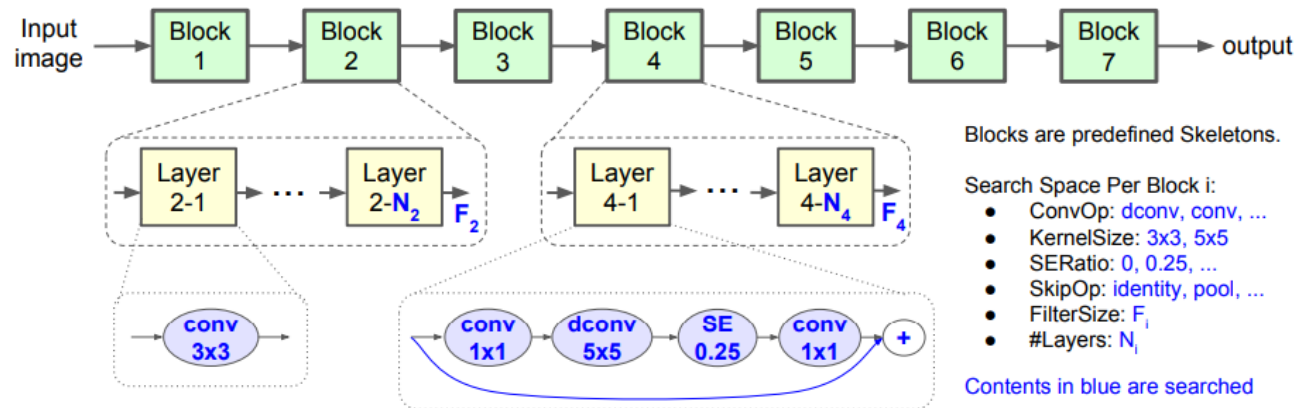
↓

$$\text{maximize}_m \quad ACC(m) \times \left[ \frac{LAT(m)}{T} \right]^w$$
$$w = \begin{cases} \alpha, & \text{if } LAT(m) \leq T \\ \beta, & \text{otherwise} \end{cases}$$



# MnasNet (Tan et al., 2019)

- Use Factorized Hierarchical Search Space
- Network layers are grouped into a number of blocks
- Each block contains a variable number of repeated identical layers
- For each block, search for the operations for a single layer and the number of layers






- Convolutional ops *ConvOp*: regular conv (conv), depthwise conv (dconv), and mobile inverted bottleneck conv [29].
- Convolutional kernel size *KernelSize*: 3x3, 5x5.
- Squeeze-and-excitation [13] ratio *SERatio*: 0, 0.25.
- Skip ops *SkipOp*: pooling, identity residual, or no skip.
- Output filter size  $F_i$ .
- Number of layers per block  $N_i$ .

# EfficientNet Architecture

- Since model scaling does not change layer, having a good baseline is critical
- Use  $ACC(m) \times [\frac{FLOPS(m)}{T}]^w$  as the optimization goal where  $w = -0.07$

**Table 1. EfficientNet-B0 baseline network** – Each row describes a stage  $i$  with  $\hat{L}_i$  layers, with input resolution  $\langle \hat{H}_i, \hat{W}_i \rangle$  and output channels  $\hat{C}_i$ . Notations are adopted from equation 2.

Stage $i$	Operator $\hat{\mathcal{F}}_i$	Resolution $\hat{H}_i \times \hat{W}_i$	#Channels $\hat{C}_i$	#Layers $\hat{L}_i$
1	Conv3x3	$224 \times 224$	32	1
2	MBConv1, k3x3	$112 \times 112$	16	1
3	MBConv6, k3x3	$112 \times 112$	24	2
4	MBConv6, k5x5	$56 \times 56$	40	2
5	MBConv6, k3x3	$28 \times 28$	80	3
6	MBConv6, k5x5	$14 \times 14$	112	3
7	MBConv6, k5x5	$14 \times 14$	192	4
8	MBConv6, k3x3	$7 \times 7$	320	1
9	Conv1x1 & Pooling & FC	$7 \times 7$	1280	1



# EfficientNet Architecture

- Starting from the baseline EfficientNet-B0, apply compound scaling method:
  - STEP1: Fix  $\phi = 1$ , assuming x2 resources. Do a grid search to find  $\alpha, \beta, \gamma$
  - STEP2: Fix  $\alpha, \beta, \gamma$  as constants and scale up baseline network with different  $\phi$
- Resultant  $\alpha = 1.2, \beta = 1.1, \gamma = 1.15$

$$\max_{d,w,r} \text{Accuracy}(\mathcal{N}(d, w, r))$$

$$s.t. \quad \mathcal{N}(d, w, r) = \bigodot_{i=1 \dots s} \hat{\mathcal{F}}_i^{d \cdot \hat{L}_i} (X_{\langle r \cdot \hat{H}_i, r \cdot \hat{W}_i, w \cdot \hat{C}_i \rangle})$$

$$\text{Memory}(\mathcal{N}) \leq \text{target\_memory}$$

$$\text{FLOPS}(\mathcal{N}) \leq \text{target\_flops}$$

$$\text{depth: } d = \alpha^\phi$$



$$\text{width: } w = \beta^\phi$$

$$\text{resolution: } r = \gamma^\phi$$

$$s.t. \quad \alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

$$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$$

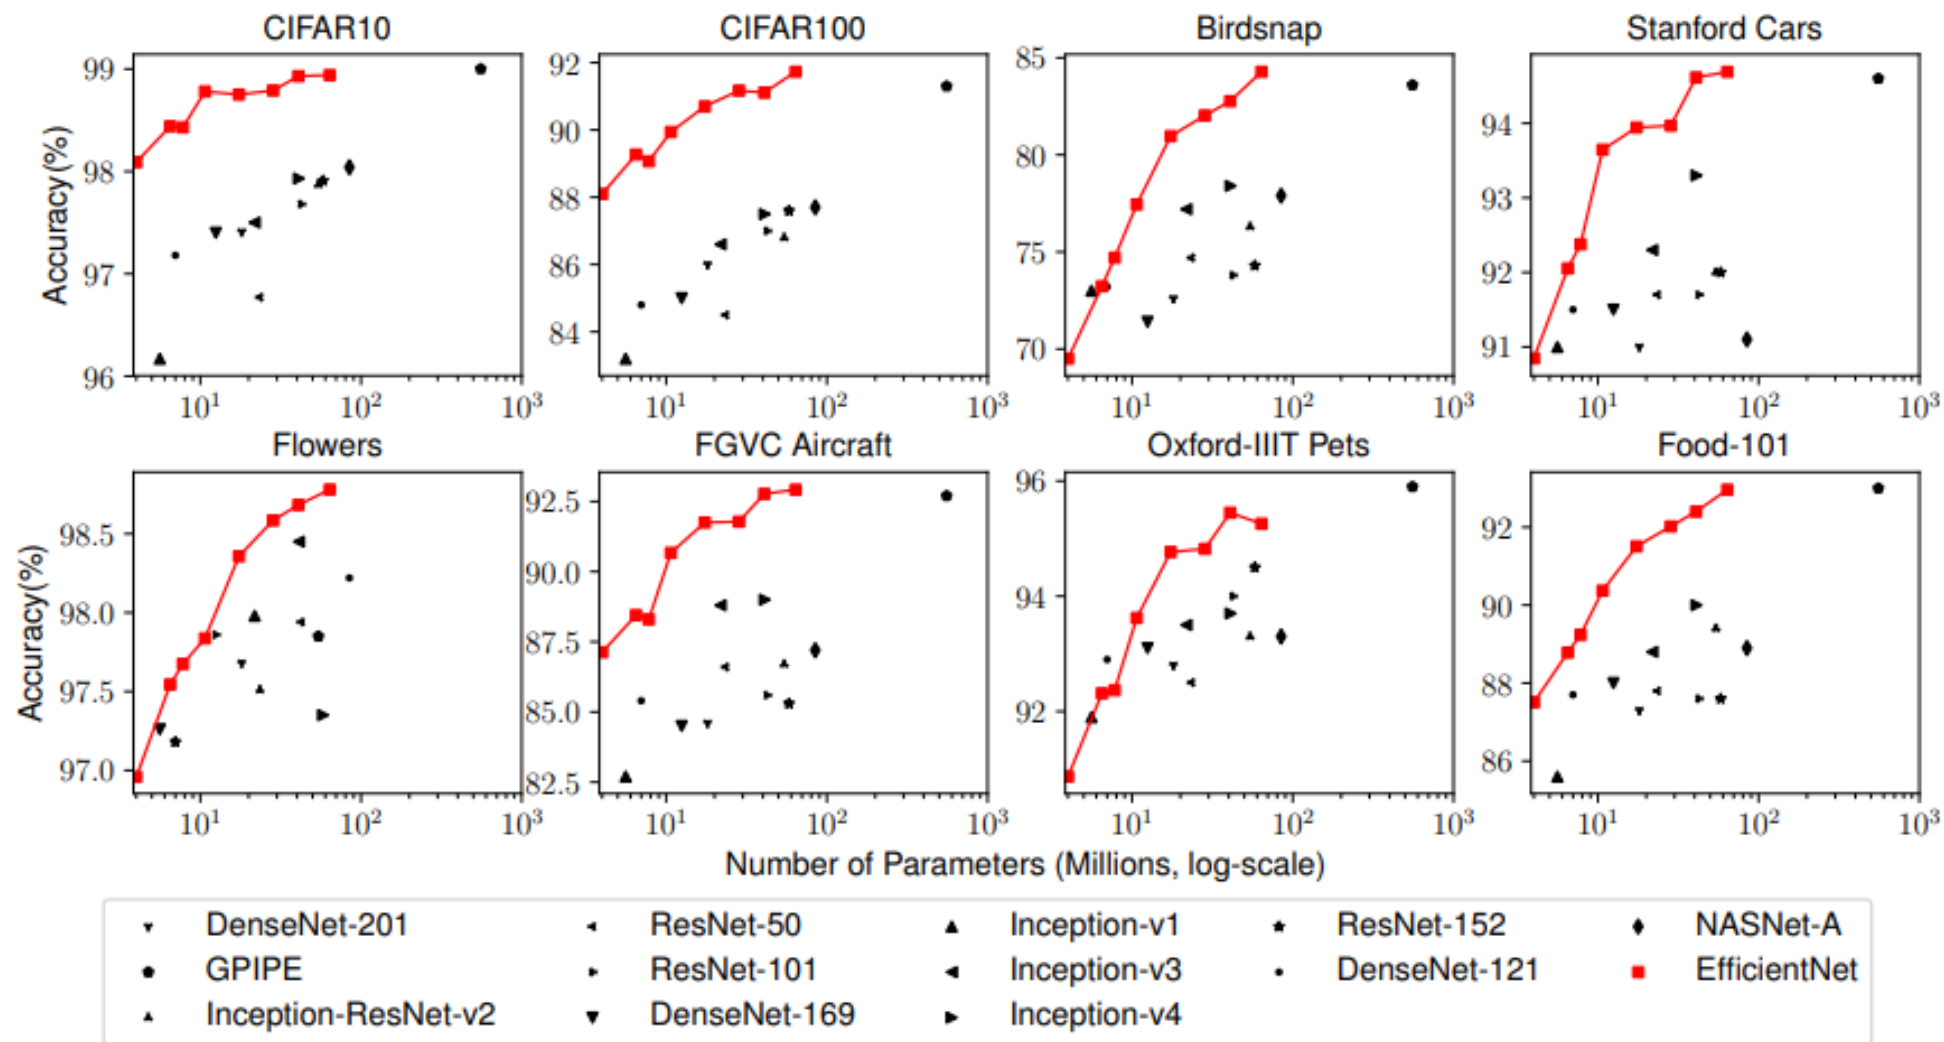
# Results on ImageNet



Model	Top-1 Acc.	Top-5 Acc.	#Params	Ratio-to-EfficientNet	#FLOPS	Ratio-to-EfficientNet
<b>EfficientNet-B0</b>	<b>77.3%</b>	<b>93.5%</b>	<b>5.3M</b>	<b>1x</b>	<b>0.39B</b>	<b>1x</b>
ResNet-50 (He et al., 2016)	76.0%	93.0%	26M	4.9x	4.1B	11x
DenseNet-169 (Huang et al., 2017)	76.2%	93.2%	14M	2.6x	3.5B	8.9x
<b>EfficientNet-B1</b>	<b>79.2%</b>	<b>94.5%</b>	<b>7.8M</b>	<b>1x</b>	<b>0.70B</b>	<b>1x</b>
ResNet-152 (He et al., 2016)	77.8%	93.8%	60M	7.6x	11B	16x
DenseNet-264 (Huang et al., 2017)	77.9%	93.9%	34M	4.3x	6.0B	8.6x
Inception-v3 (Szegedy et al., 2016)	78.8%	94.4%	24M	3.0x	5.7B	8.1x
Xception (Chollet, 2017)	79.0%	94.5%	23M	3.0x	8.4B	12x
<b>EfficientNet-B2</b>	<b>80.3%</b>	<b>95.0%</b>	<b>9.2M</b>	<b>1x</b>	<b>1.0B</b>	<b>1x</b>
Inception-v4 (Szegedy et al., 2017)	80.0%	95.0%	48M	5.2x	13B	13x
Inception-resnet-v2 (Szegedy et al., 2017)	80.1%	95.1%	56M	6.1x	13B	13x
<b>EfficientNet-B3</b>	<b>81.7%</b>	<b>95.6%</b>	<b>12M</b>	<b>1x</b>	<b>1.8B</b>	<b>1x</b>
ResNeXt-101 (Xie et al., 2017)	80.9%	95.6%	84M	7.0x	32B	18x
PolyNet (Zhang et al., 2017)	81.3%	95.8%	92M	7.7x	35B	19x
<b>EfficientNet-B4</b>	<b>83.0%</b>	<b>96.3%</b>	<b>19M</b>	<b>1x</b>	<b>4.2B</b>	<b>1x</b>
SENet (Hu et al., 2018)	82.7%	96.2%	146M	7.7x	42B	10x
NASNet-A (Zoph et al., 2018)	82.7%	96.2%	89M	4.7x	24B	5.7x
AmoebaNet-A (Real et al., 2019)	82.8%	96.1%	87M	4.6x	23B	5.5x
PNASNet (Liu et al., 2018)	82.9%	96.2%	86M	4.5x	23B	6.0x
<b>EfficientNet-B5</b>	<b>83.7%</b>	<b>96.7%</b>	<b>30M</b>	<b>1x</b>	<b>9.9B</b>	<b>1x</b>
AmoebaNet-C (Cubuk et al., 2019)	83.5%	96.5%	155M	5.2x	41B	4.1x
<b>EfficientNet-B6</b>	<b>84.2%</b>	<b>96.8%</b>	<b>43M</b>	<b>1x</b>	<b>19B</b>	<b>1x</b>
<b>EfficientNet-B7</b>	<b>84.4%</b>	<b>97.1%</b>	<b>66M</b>	<b>1x</b>	<b>37B</b>	<b>1x</b>
GPipe (Huang et al., 2018)	84.3%	97.0%	557M	8.4x	-	-

We omit ensemble and multi-crop models (Hu et al., 2018), or models pretrained on 3.5B Instagram images (Mahajan et al., 2018).

# Results on Transfer Learning





# Comparisons with different scaling methods

- The model with compound scaling tends to focus more relevant regions

