

Scalpel: Customizing DNN Pruning to the Underlying Hardware Parallelism

ISCA 17

Neural Network Quantization & Compact Network Design Study, AI Robotics KR

Constant (Sang-Soo) Park

<https://constantpark.github.io/>

2019-10-20

Contents


- Introduction
- Background & Motivations
- Contributions
- Scalpel: SIMD-aware weight/Node pruning
- Evaluation
- Conclusion

Introduction

- **Pruning deep neural networks to make them fast and small**
 - Removing connections that don't affect significantly affect the results^[1]
 - Possible to reduce the number parameters, the amount of computation
 - **But will it really run faster in modern processor and dedicated processor?**

Table 1: The compression pipeline can save $35\times$ to $49\times$ parameter storage with no loss of accuracy.

Network	Top-1 Error	Top-5 Error	Parameters	Compress Rate
LeNet-300-100 Ref	1.64%	-	1070 KB	
LeNet-300-100 Compressed	1.58%	-	27 KB	40\times
LeNet-5 Ref	0.80%	-	1720 KB	
LeNet-5 Compressed	0.74%	-	44 KB	39\times
AlexNet Ref	42.78%	19.73%	240 MB	
AlexNet Compressed	42.78%	19.70%	6.9 MB	35\times
VGG-16 Ref	31.50%	11.32%	552 MB	
VGG-16 Compressed	31.17%	10.91%	11.3 MB	49\times



Background: Pruning

- **Fine-grained vs Course-grained pruning**
 - Fine-grained: Random distribution of zero values^[2]
 - Course-grained: **Structural pruning** such as channel-pruning^[3]

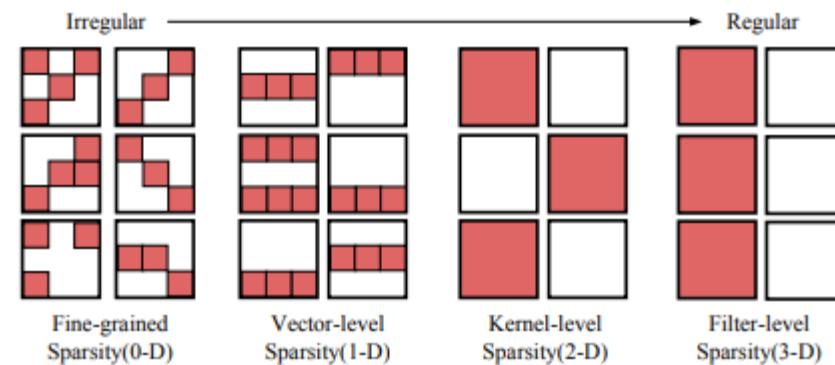
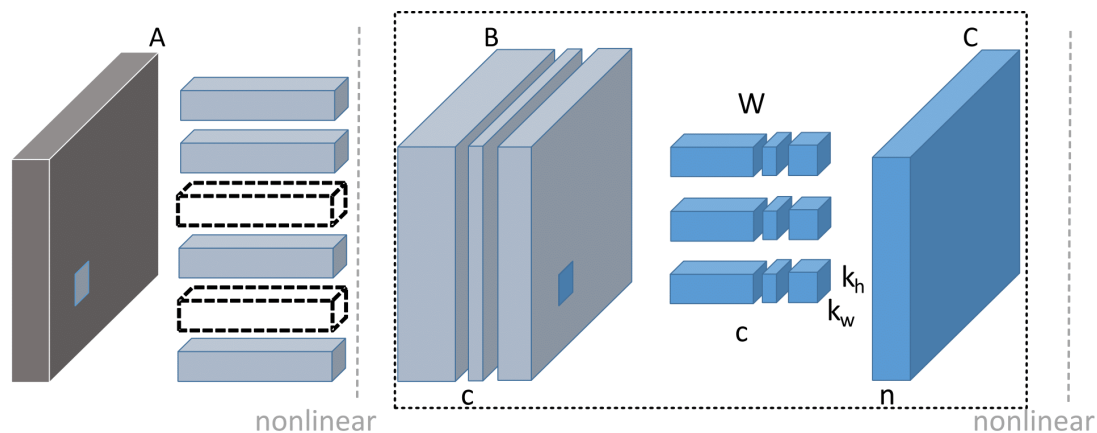


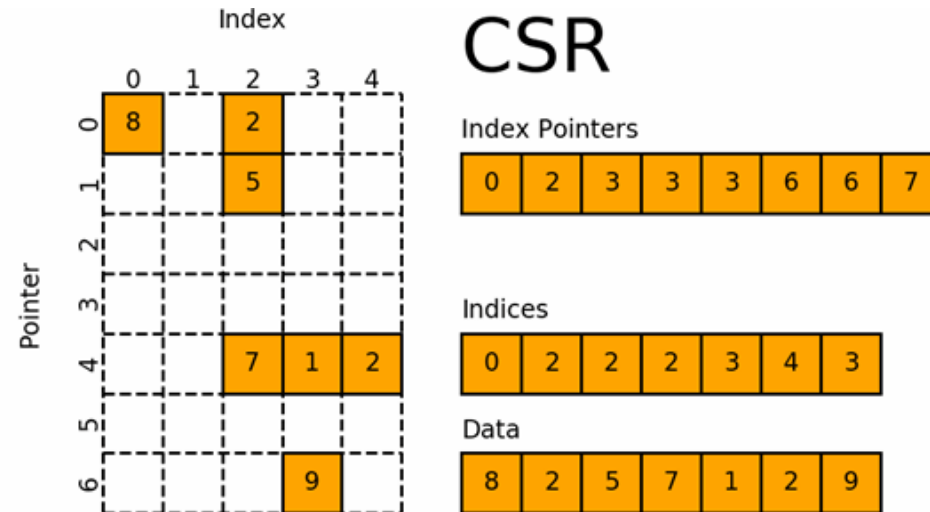
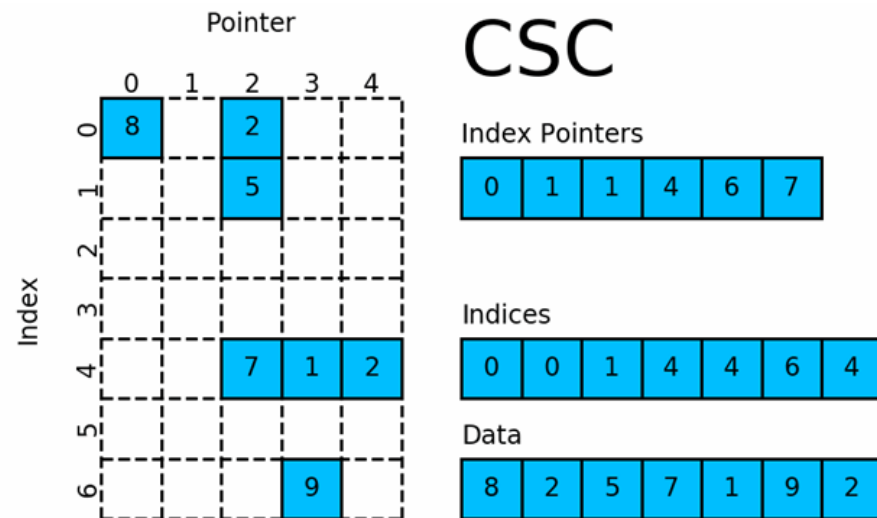
Figure 1: Different sparse structure in a 4-dimensional weight tensor. Regular sparsity makes hardware acceleration easier.

[2] <https://github.com/yihui-he/channel-pruning>

[3] [https://github.com/NervanaSystems/distiller/wiki/Frequently-Asked-Questions-\(FAQ\)](https://github.com/NervanaSystems/distiller/wiki/Frequently-Asked-Questions-(FAQ))

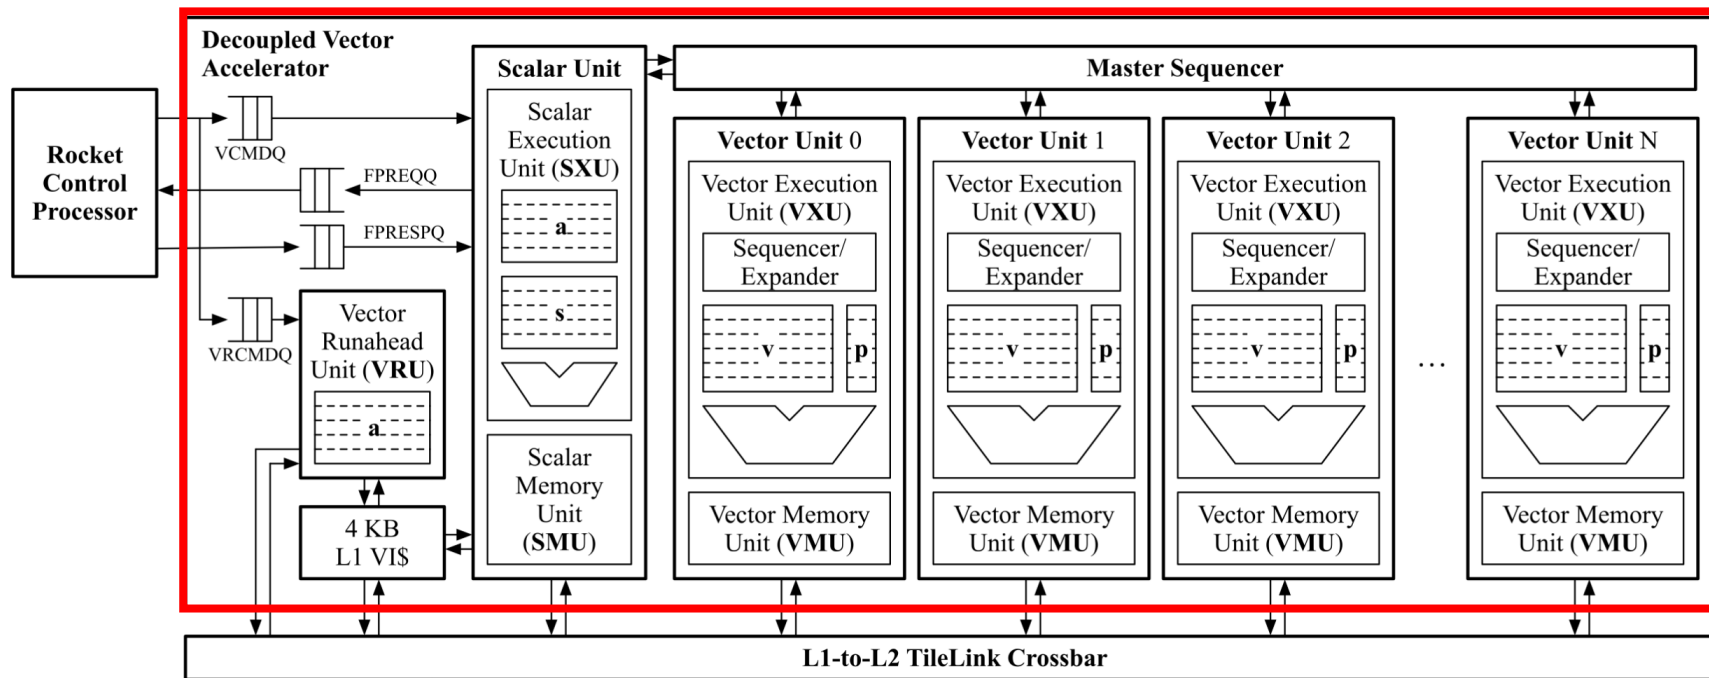
Background: Sparse matrix format

- **Storage format for storing matrix**
 - Store non-zero values and **index information (where this value is stored)**
 - Many ways to store sparse matrix: COO, CSR, CRS, etc.
 - Require additional decoding computations




Background: Vector processor (SIMD)

- **Single instruction multiple data (SIMD)**
 - Dedicated accelerator for vector processing inside processor
 - Compute arithmetic operations **in parallel** using **SIMD accelerator**
 - **# of Lane: # of arithmetic operations that can be computed in parallel**



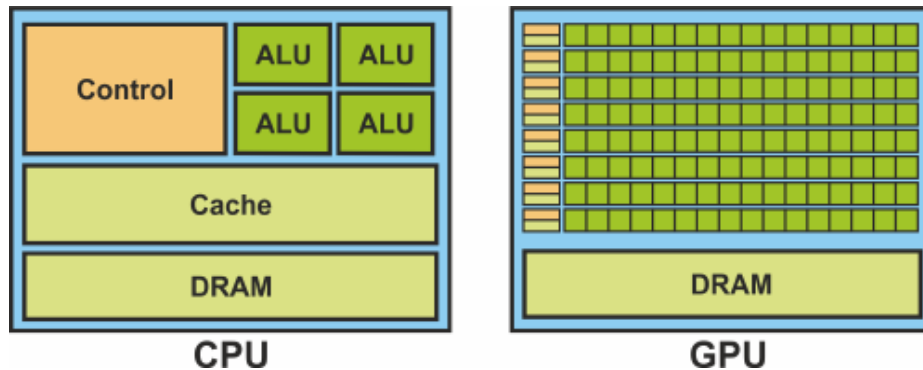
Motivation: Vector processor (SIMD)

- **Problems in SIMD operations**
 - In the case of **sparse matrix-vector multiplication (SPMV)**
 - **Maximum vector length (# of lane) is 4**
 - Vector length of 3: 3 computations in parallel, but one lane is idle




Background: Vector processor (SIMT)


- **Single instruction multiple thread (SIMT)**
 - Graphic processing unit (GPU): **many core architecture**
 - CPU: 1~16 cores, GPU: more than 1,000 cores
 - Threads are grouped and controlled by group of thread (thread-block/work-group, warp/wave-front)



Pre-Volta



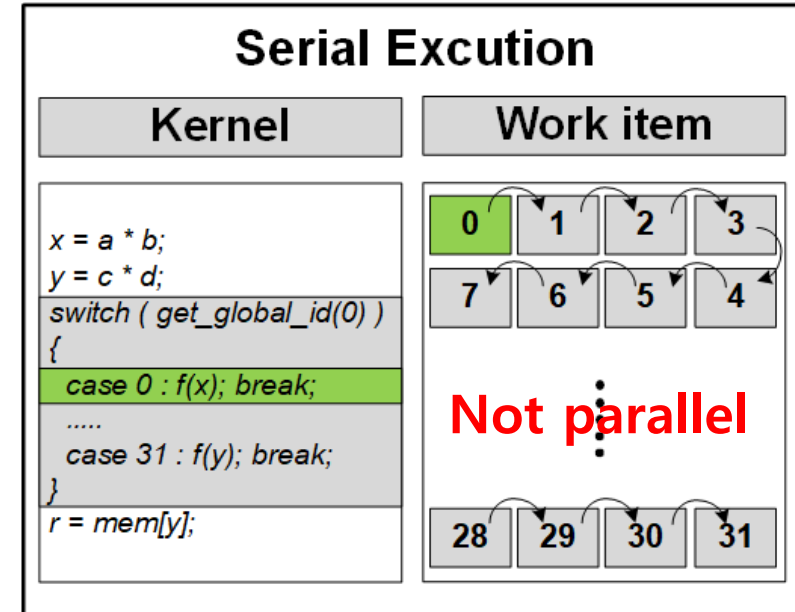
Volta



Motivation: Vector processor (SIMT)


- **Divergence branch in GPU**
 - Conditions (If-else, Case): serial execution by branch instruction, Idle state computing units (PEs)^[4]
 - **Decoding makes additional condition loops**

```
__global__ void SpMV_CSR_kernel(const SparseMatrixCSR A, const float * x, float * y) {  
  
    const int row = blockIdx.x * blockDim.x + threadIdx.x;  
  
    if (row < A.M) {  
  
        float dotProduct = 0;  
        const int row_start = A.row_indices[row];  
        const int row_end = A.row_indices[row+1];  
        for (int element = row_start; element < row_end; ++element) {  
  
            dotProduct += A.values[element] * x[A.col_indices[element]];  
  
        }  
  
        y[row] = dotProduct;  
  
    }  
  
}
```



Contributions

- DNN weight pruning is not a panacea
 - Performance affected by both the structure network and data-parallel hardware
 - Pruned network that is customized to hardware platform (CPU, GPU)
 - Depend on hardware parallelism and layer type, choose **SIMD-aware weight pruning** or **Node-pruning**



Scalpel: Pruning technique

- Customized DNN pruning for different hardware platform
 - Two pruning techniques: SIMD-aware weight punning/Node pruning

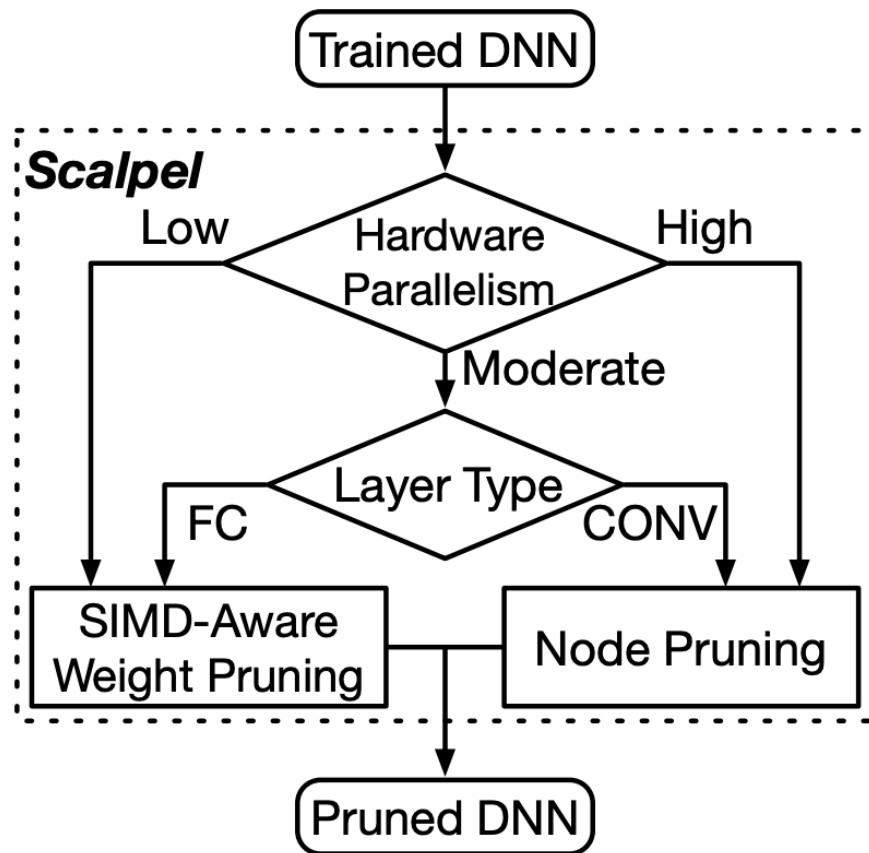


Table 1: Hardware platforms with different parallelism.

	Parallelism		
	Low	Moderate	High
Example	Micro-controller	CPU	GPU
Memory Hierarchy	No cache	Deep cache hierarchy	High bandwidth / long latency
Memory Size	~100KB SRAM	~8MB SRAM	2-12GB DRAM

Scalpel: SIMD-aware weight pruning

- **Customized DNN pruning for different hardware platform**
 - Grouping weights by the number of SIMD lanes
 - Using root-mean-square (RMS), measure the importance of weight groups
 - Removing grouped weights under threshold and retraining
 - Layer by layer manner with dropout

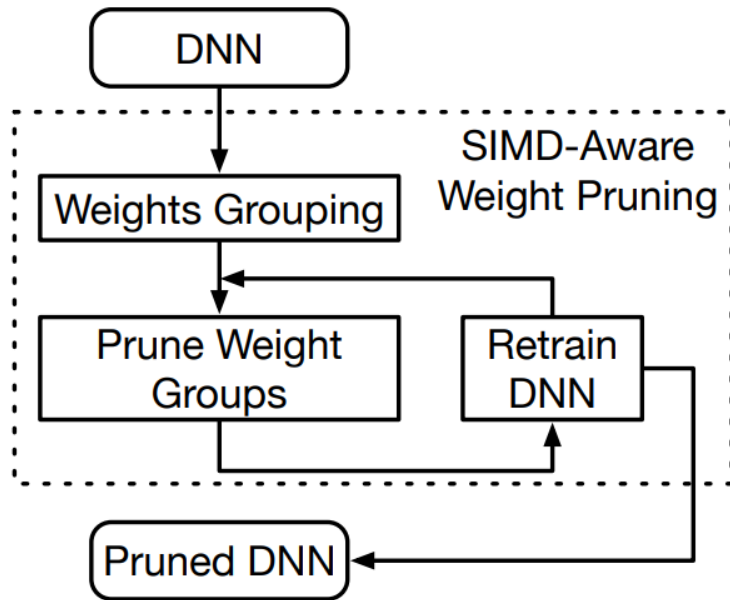


Figure 7: Main steps of SIMD-aware weight pruning.




Figure 8: (A) Weights grouping; (B) Sparse weight matrix after pruning weight groups; (C) Modified CSR format for SIMD-aware weight pruning.

Scalpel: Node pruning (=Channel pruning)

- **Hardware friendly pruning technique**

- Sparse matrix multiplication requires more execution time compared to dense matrix
- Using **mask layer** (Boolean α): blocking one node (one neuron in FC, one feature map in convolution layer)
- L1 regularization to control the number of nodes got removed
- If the node is not important, β will be decreased to be lower than threshold T

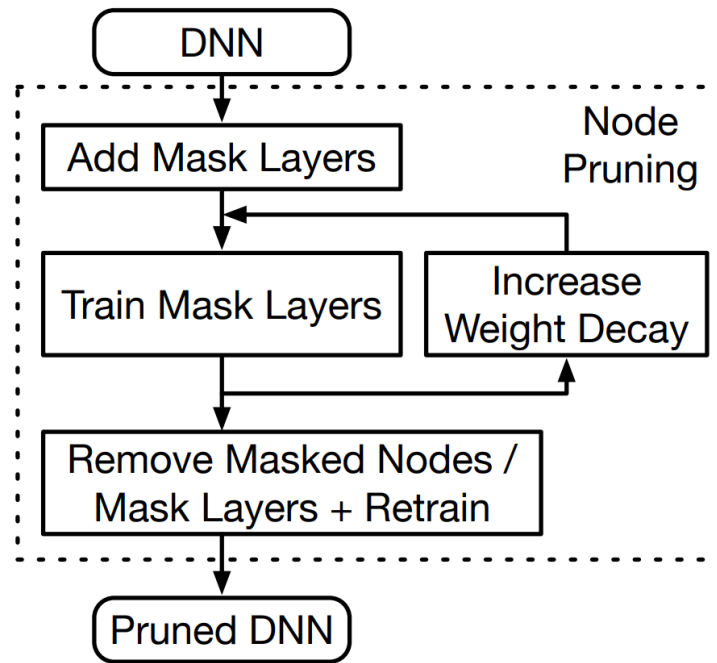


Figure 11: Main steps of node pruning.

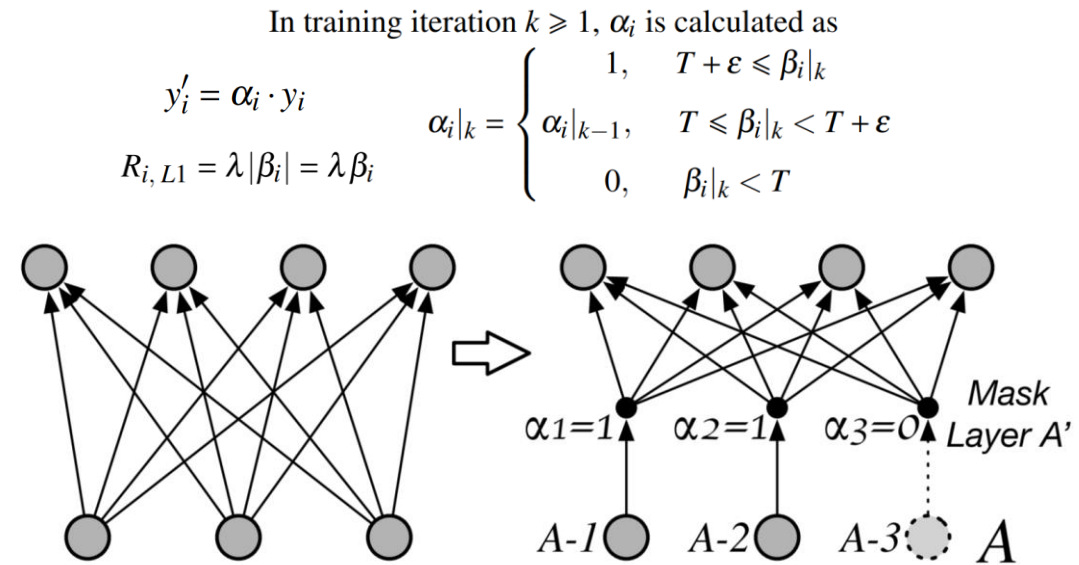


Figure 12: Mask layers. Node A-3 with $\alpha_3 = 0$ can be removed. The whole mask layer A' will be removed after pruning all redundant nodes.

Evaluation: Hardware platform

- **Micro-controller unit (MCU): low parallelism**
 - ARM Cortex-M4: 3-stage in-order pipeline
 - **2-way SIMD unit** for 16b fixed-point
 - ARM's sparse matrix-vector BLAS
- **CPU: moderate parallelism**
 - Intel Core i7-6700 Skylake
 - 8-way SIMD ISA for 32b float
 - Math kernel library (MKL) CSRMMV/CSRMM
- **GPU: high parallelism**
 - NVIDIA Titan X
 - cuSPARSAE CSRMMV/CSRMM

Evaluation: Speedup in CPU

- **Speedup in i7-6700 CPU running convolution and fully-connected**
 - SIMD-Sparse (SIMD-aware), Simple-Sparse (Sparse matrix), Intel's BLAS (MKL-Dense/MKL-Sparse)
 - **Conventional sparse methods** (MKL/Simple-Sparse) require more execution time compared to dense
 - **To achieve performance speedup with pruning, at least 79% of weights need to be removed**
 - But conv layers have much less redundancy compared fully-connected layer




Figure 14: Relative execution time for sparse matrix-matrix multiplication (CONV layers) on Intel Core i7-6700. The weight matrix and input matrix have the size of 128 x 1200 and 1200 x 729, respectively.




Figure 13: Relative execution time for sparse matrix-vector multiplication (FC layers) on Intel Core i7-6700. The matrix size is 4096 x 4096 and the vector size is 4096. MKL-Dense/Sparse show the results of dense and sparse weight matrix with the Intel MKL library.

Evaluation: Model size in CPU

- **Improvement of speedup and model size**
 - Speedup can be up to **6.86x**, relative size is reduced by up to **94.8%**
 - **Non of the layer in NIN can have more than 50% of weights**
 - **Converting to sparse matrix is expected to require additional overhead.**

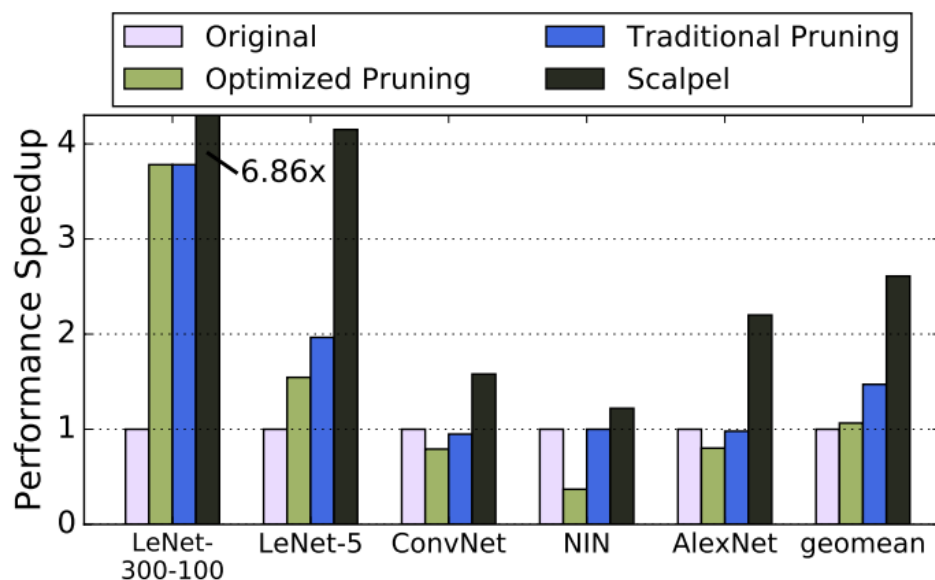


Figure 18: Relative performance speedups of the original models, traditional pruning, optimized pruning and Scalpel on Intel Core i7-6700 CPU.




Figure 19: Relative model sizes of the original models, traditional pruning, optimized pruning and Scalpel for Intel Core i7-6700 CPU.

Evaluation: High parallelism in GPU

- Improvement of speedup and model size
 - LeNet-300-100 is very **tiny** model
 - Removing DNN redundancy at a granularity of nodes to keep regular structure**

DNNs	Percentage of Nodes Removed in Each Layer
LeNet-300-100	31% (fc1)- 32% (fc2)
LeNet-5	50% (conv1)- 68% (conv2)- 65% (fc3)
ConvNet	28% (conv1)- 25% (conv2)- 49% (conv3)
NIN	28% (conv1)- 20% (cccp1)- 5% (cccp2)- 2% (conv2)- 14% (cccp3)- 8% (cccp4)- 22% (conv3)- 48% (cccp5)
AlexNet	3% (conv1)- 20% (conv2)- 24% (conv3)- 18%(conv4)-0%(conv5)-17%(fc6)-23%(fc7)




Figure 20: Relative performance speedups of the original models, traditional pruning, optimized pruning and Scalpel on NVIDIA GTX Titan X GPU.




Figure 21: Relative model sizes of the original models, traditional pruning, optimized pruning and Scalpel for NVIDIA GTX Titan X GPU.

Evaluation: Compared to STOA

- **Performance close to Deep Compression^[1]**
 - **Deep Compression:** Pruning + Quantization + Lookup table
 - **Scalpel:** 32b floating-point + Pruning
 - In the case of AlexNet
 - Model size: Deep Compression (13.06%), Scalpel (11%)
 - Speedup: Deep Compression (3x), Scalpel (2.2x)

In Deep Compression	
DNNs	Relative Size
LeNet-300-100	8%
LeNet-5	8%
AlexNet	11%
VGG-16	7.50%

Table 3: Results overview.

Hardware	DNNs	Speedup	Relative Size
Micor-controller	LeNet-300-100	9.17x	6.93%
	LeNet-5	3.51x	6.72%
	ConvNet	1.38x	40.95%
CPU	LeNet-300-100	6.86x	7.08%
	LeNet-5	4.15x	5.20%
	ConvNet	1.58x	44.28%
	NIN	1.22x	81.16%
	AlexNet	2.20x	13.06%
GPU	LeNet-300-100	1.08x	66.83%
	LeNet-5	1.59x	11.67%
	ConvNet	1.14x	45.40%
	NIN	1.17x	81.16%
	AlexNet	1.35x	76.52%




Figure 9: Compared with the original network, pruned network layer achieved 3× speedup on CPU, 3.5× on GPU and 4.2× on mobile GPU on average. Batch size = 1 targeting real time processing. Performance number normalized to CPU.




Figure 10: Compared with the original network, pruned network layer takes 7× less energy on CPU, 3.3× less on GPU and 4.2× less on mobile GPU on average. Batch size = 1 targeting real time processing. Energy number normalized to CPU.

Conclusion

- **Scalpel to customize DNN pruning for different hardware platform**
 - SIMD-aware weight pruning, Node pruning
 - On MCU, CPU, and GPU
 - 3.54x, 2.61x, and 1.25x of speedup
 - 88%, 82%, 53% of reduced model size
- **Potential problem**
 - Activation pruning when **zero** value

